



DSA211 Statistical Learning with R

AY 23/24 Term 2

Section: G1

Professor: Kwong Koon Shing

Group Project

Group Members	Student ID
Ariel Wong Wei Lin	01465744
Tan Yu Xuan	01408987
Yip Yu Liang Adrian	01425210
Quek Jia Yang	01373129
Reese Lim Kaylee	01472009

Table of Contents

Prologue	3
Models and Analyses	3
1. Multiple Linear and Polynomial Regression	3
2. Best Subset Selection	4
3. Ridge Regression	4
4. The Lasso	4
5. Elastic Net Regression	5
6. K-nearest Neighbours Algorithm	5
7. Decision Tree	5
8. Random Forest Ensemble method	6
Final Model Selection	7
Appendix	8
Appendix 1 (Multiple linear and polynomial regression)	8
Appendix 2 (Best Subset Selection)	12
Appendix 3 (Ridge Regression)	15
Appendix 4 (The Lasso)	16
Appendix 5 (Elastic Net Regression)	18
Appendix 6 (K-nearest Neighbours)	20
Appendix 7 (Decision Tree)	21
Appendix 8 (Random Forest)	27

Prologue

We used the summary function on the Central dataset first and identified that *Tenure*, *Purchaser* and *Region* are categorical variables while the rest are numerical.

```
> summary(cen)
      Price.V1      Area      Age      Tenure
Min.   :-1.021308 Min.   : 398.3 Min.   : 0.00 Freehold : 914
1st Qu.: -0.563140 1st Qu.: 818.1 1st Qu.: 7.00 Leasehold:1586
Median : -0.261263 Median :1141.0 Median :12.00
Mean    : 0.000000 Mean    :1279.0 Mean    :13.54
3rd Qu.: 0.218202 3rd Qu.:1517.7 3rd Qu.:19.00
Max.    :10.979402 Max.    :7717.8 Max.    :45.00

Purchaser      Region
HDB           : 626  Bukit Timah:427
Private:1874   Bukit Merah:306
               Toa Payoh :301
               Kallang   :217
               Bishan    :198
               Queenstown:173
               (Other)   :878
```

We construct all models on a sample size of 2000 for training, and use the remaining 500 to calculate Mean Squared Error (MSE) of the test set.

Models and Analyses

1. Multiple Linear and Polynomial Regression

Firstly, we run a multiple linear regression taking all variables into account using the `lm` function. We note that *Area*, *Age*, *Tenure*, and selected *Region* dummy variables are significant with a p-value less than 0.01, while *Purchaser* has a p-value of 0.9256. This seems to suggest that *Purchaser* is a less significant predictor variable in determining a change in *Price*.

We then run through multiple combinations of variables, by varying the terms to include inside interaction terms, and also varying the polynomials of *Area* and *Age*. Furthermore, for the sake of simplicity, we do not include *Region* inside the interaction terms at all. For example, we ran models with interaction terms of *Tenure* and *Area*, *Area* and *Age*, and *Age* and *Purchaser*. We also determined the best polynomial for *Area* and *Age* by running LOOCV on each of them by itself, and determined the best was $Area^3$ and Age^2 . For a fairer comparison, we will compare the adjusted R-squared across every model to identify the best fitting model.

Overall, we identified that a multiple regression with the interaction terms of *Area* and *Tenure* is the best, giving rise to an adjusted R-squared of 0.8499. We also calculated the MSE to be 481090502793 for comparison between other models. The coefficients of the best model is as given in Appendix 1.

2. Best Subset Selection

We tried running a best subset selection with 10-fold cross-validation, where we choose a subset of k variables out of p that gives us the best model in terms of smallest MSE. Through cross-validation, we found the model with the smallest MSE of 507174939759 when $k = 12$ (Appendix 2). Best subset selection involves evaluating all possible combinations of variables (out of 18 variables), and when combined with cross-validation, its computational complexity is very high, since the process needs to be repeated for each fold of the cross-validation. Furthermore, with best subset selection, there is a risk of overfitting the model to the training data, especially when the number of variables is large compared to the number of observations. Finally, its MSE is worse than the initial polynomial multiple regression model. Hence, we do not choose to use the best subset selection for this problem.

3. Ridge Regression

We want to carry out a shrinkage approach to see if Ridge regression could be used to fit all predictors and constrain the coefficient estimates towards zero, to prevent overfitting. The first method we will use for this is ridge regression. Under linear regression, we estimate coefficients based on values that minimise residual sum of squares. Ridge regression also does this, but has an extra penalty term which is determined by the sum of squares of coefficients multiplied by a tuning parameter λ . The penalty term allows us to take multicollinearity into account as it introduces bias to the model - as λ increases, bias increases while variance decreases. We carry out cross validation to determine the best value of λ with the lowest cross validation error. We end up with a λ of 158285. This λ then produces a model with an MSE of 519472000000 and coefficient estimates as shown in the figure in Appendix 3. This is worse than the best subset selection MSE = 507174939759. Additionally, the drawback of the ridge regression is that it will include all predictors in the final model. Hence, we also carried out the Lasso.

4. The Lasso

The Lasso can prevent overfitting while overcoming the disadvantage of ridge regression, as it also implements variable selection by shrinking coefficients to exactly 0. Similar to the ridge regression, the Lasso estimates coefficients based on values that minimise residual sum of squares. It also introduces a penalty term based on the absolute values of the coefficients multiplied by a tuning parameter λ . The penalty term allows the Lasso regression to shrink the coefficients towards 0. When λ is sufficiently large, coefficients can shrink to exactly 0, giving rise to variable selection. Likewise, cross validation is carried out to determine the best value of λ with the lowest cross validation error. We get a λ value of 1778.053 which produces a model with an MSE of 506696419155 and coefficient estimates as shown in Appendix 4. This MSE is lower than that of the model produced by ridge regression, and *Purchaser* is dropped in the final model of the Lasso, verifying its relative insignificance as suggested in the first proposed model.

5. Elastic Net Regression

Following ridge regression and the Lasso, we also decided to try a combination of both regularisation methods by using elastic net regression, which is useful when dealing with multicollinearity and overfitting issues. Given that $\alpha=0$ is set for ridge regression and $\alpha=1$ is set for Lasso, we initially experimented with elastic net regression by taking the middle value and setting $\alpha=0.5$, which produces an MSE of 506676319402 with best value of λ of 3240.192.

However, since there are different types of elastic net regressions — some being lasso-dominant and others more ridge-dominant — we explored further and considered other cases of $0 < \alpha < 1$, where α need not necessarily be in the middle. Therefore, we considered 2 other α values, $\alpha=0.75$ (Lasso-dominant) and $\alpha=0.25$ (Ridge-dominant). When $\alpha=0.75$, it returns an MSE of 506795195894 with a best value of λ of 2370.737. On the other hand, when $\alpha=0.25$, the model returns an MSE value of 506447279753 and best λ value of 5380.128.

Our findings show that this project is more suited to adopt a mixture of ridge regression and the Lasso with a more ridge-dominant approach, since the elastic net regression model performs best when α is set to 0.25 as it gives the smallest MSE. Appendix 5 shows the final model of elastic net regression, which excludes the *Purchaser* variable.

6. K-nearest Neighbours Algorithm

We then attempted a K-nearest neighbour (KNN) approach, which is a non-parametric method used for both classification & regression. KNN consists of using a distance measure (Euclidean distance, Manhattan, Mikowski & Hamming) to group objects into classes. In our case, we used Euclidean distance based on restrictions of R's KNN package. We chose $K = 4$ after testing, as higher values increase the MSE value. In regression, the predicted *Price* of the test data set is calculated as the average of the value of the 4 nearest data points in a class. Unlike other methods, KNN is susceptible to large differences in scale of data as, in our case *Area* & *Age*, where *Area* is especially high. This will cause unequal weight to be placed upon these variables during distance calculation, and choose to scale all predictor columns. The result is that even after scaling data for KNN & comparing it to scaled calculated MSE's of our other models, the $MSE = 8800808162478.64$ which was still the worst of all of them. The unscaled $MSE = 9285272687764.77$, which is still the worst. We decided not to pursue this method further. The scaled value can be obtained by uncommenting `x[,2:6] <- scale(x[,2:6])`.

7. Decision Tree

Given the poor performance of KNN, we decided to apply tree based methods as they are suitable for both regression & classification. Despite these being inferior in accuracy to supervised learning approaches, we believe a decision tree would be suitable for price

predictions, as it is able to handle non-linear & interaction effects, which based on earlier analysis of the linear model, show the highest adjusted R-squared with an interaction effect between *Area* & *Tenure*.

In addition, decision trees are able to handle the mixed data we have in our dataset as *Region* *Purchase* & *Tenure* are categorical as opposed to the other columns which are continuous. As before, we chose the scale of the *Price* column to have more readable MSE values.

In our first construction of the tree, there were 8 terminal nodes, as seen in Appendix 7, “Regression Tree for Central2024P data”.

However, we need to consider the effect of an excessively large number of terminal nodes which generates complexity when fitting the data. This increases MSE and reduces test accuracy beyond a certain number of nodes. We will adopt cross-validation to prune the tree by identifying the largest number of terminal nodes, with the lowest deviance. Based on our codings, having 7 terminal nodes is the best as it gives the lowest deviance of 1.552796×10^{15} . This is better in comparison to having 8 terminal nodes which gives a deviance of 1.565915×10^{15} . See Appendix 7 Cross Validation: Deviance versus Size.

Given this, we now prune the original decision tree using the prune functions within R, and predict the *Price* with this pruned tree to the actual test prices. The result is an MSE of 6.30826×10^{11} , with the tree seen in Appendix 7 “Pruned Regression Tree for Central Data”, and the predictions in Appendix 7 “Pruned Tree predictions versus observed prices for test data”.

We now rebuild a tree on the full dataset, with the constraint of maximum terminal nodes to 7, and predict its outputs, with comparison to the test dataset. The final tree can be referred to in appendix 7 “Pruned Regression Tree for all Central Data”. The MSE = 625661162260, indicating the restriction to a maximum of 7 terminal nodes is ideal. This result is worse than aforementioned models such as multiple polynomial regression, Ridge, Lasso & Elastic Net Regression.

8. Random Forest Ensemble method

Several flaws exist with using decision trees, namely that it runs the risk of overfitting, and is sensitive to changes in training data, resulting in high variance. It is also a greedy algorithm and cannot guarantee the optimal tree.

Our solution is to use the Random Forest (RF) ensemble method. RF is a type of bagging method that involves building multiple decision trees on subsets of data randomly sampled with replacement from the train dataset. Multiple different trees are thus created, trained on a different subset of the training dataset. The final prediction is obtained by averaging predictions across the multiple decision trees created for the final prediction. This has several advantages compared to a single decision tree. Its use of aggregates from multiple trees prevents overfitting, so a more robust model is trained. Additionally, the usage of multiple

trees accounts for multiple combinations, simulating new data, meaning it can generalise well to completely novel datasets. Finally, RF ensemble methods work well on high-dimension data, hence it can handle the high number of features we have arising due to the multiple locations we have in the *Region* variable.

In executing the random forest ensemble method, we decided to modify the number of trees and eventually choose 550 trees as increasing it further will decrease our MSE, but raise the risk of overfitting even with Random Forest. In our research the ideal variable count was the square root of our total features to be estimated (18), which we rounded to 4 for variables randomly sampled at each split. In our analysis, the $MSE = 221844798581.846$ which is the lowest thus far. The diagram of the Random Forest can be seen in Appendix 8 rf_model.

%incMSE refers to the increase in MSE that comes from permuting the values of a predictor variable. The higher the %incMSE, the more important the variable in predicting an accurate response vice versa.

IncNodePurity refers to the increase in node purity when splitting on a particular variable in construction of the tree, which measures homogeneity of observations. The higher the IncNodePurity the better the split, based on that variable, leading to better performance of the random forest model.

In addition, we have also identified the 3 most important variables based on %incMSE & IncNodePurity which are *Area*, *Region*, *Age & Tenure* in order of decreasing importance. The least important variable in predicting MSE is *Purchaser* in line with our earlier analysis.

Final Model Selection

Comparing our various models, we end up with the following MSE's across all of them:

1. Multiple Linear & Polynomial Regression: $MSE = 481090502793$
2. Best Subset Selection: $MSE = 507174939759$, $k = 12$
3. Ridge Regression: $MSE = 519472000000$, $\text{Lambda} = 158285$
4. The Lasso: $MSE = 506696419155$, $\text{Lambda} = 1778.053$
5. Elastic Net Regression: $MSE = 506447279753$, $\text{Lambda} = 5380.128$ ($\alpha = 0.25$)
6. KNN: $MSE = 1202820924505.22$
7. Decision Tree: $MSE = 625661162260$
8. Random Forest: $MSE = 221844798581.846$, $\text{Trees} = 550$

Based on all our model training and evaluation, the recommended model for the dataset "Central2024P.csv" is the Random Forest ensemble method, as it has the lowest MSE of all our tested methods.

Appendix

Appendix 1 (Multiple linear and polynomial regression)

Code:

```
set.seed(9876)
cen <- read.csv("Central2024P.csv", stringsAsFactors = TRUE)
train <- sample(1:nrow(cen), 2000)
test <- (-train)
#Best Model according to Adjusted R^2
L3 <- lm(Price~Area*Tenure+Region+Age+Purchaser, cen[train,])
summary(L3)
pred3 <- predict(L3, newdata=cen[test,])
mean((pred3-cen[test, "Price"])^2)
coef(L3)

#Examples of other lm models
summary(lm(Price~.+I(Area^2)+I(Area^3), cen[train,]))
summary(lm(Price~Tenure*Purchaser+Region+Area+Age+I(Area^2), cen[train,]))
summary(lm(Price~Tenure*Purchaser+Region+Area+Age, cen[train,]))
```

Code Output:

```
> set.seed(9876)
> cen <- read.csv("Central2024P.csv", stringsAsFactors = TRUE)
> train <- sample(1:nrow(cen), 2000)
> test <- (-train)
> #Best Model according to Adjusted R^2
> L3 <- lm(Price~Area*Tenure+Region+Age+Purchaser, cen[train,])
> summary(L3)
```

Call:

```
lm(formula = Price ~ Area * Tenure + Region + Age + Purchaser,
    data = cen[train, ])
```

Residuals:

Min	1Q	Median	3Q	Max
-5421625	-177132	10246	189176	10025849

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-40570.75	87435.09	-0.464	0.642692
Area	2367.35	32.79	72.201	< 2e-16 ***
TenureLeasehold	233499.07	74386.36	3.139	0.001720 **
RegionBukit Merah	206888.80	71261.22	2.903	0.003734 **
RegionBukit Timah	77466.46	70699.90	1.096	0.273340
RegionGeylang	-70344.66	84950.66	-0.828	0.407734
RegionKallang	26012.77	78429.83	0.332	0.740174
RegionMarine Parade	284265.38	92756.38	3.065	0.002209 **

RegionNewton	1357491.90	118363.64	11.469	< 2e-16	***
RegionNovena	166281.26	95616.58	1.739	0.082183	.
RegionOthers	777659.39	114782.88	6.775	1.63e-11	***
RegionQueenstown	-49387.97	82469.45	-0.599	0.549332	
RegionRiver Valley	975827.02	102839.68	9.489	< 2e-16	***
RegionRochor	70015.97	128011.34	0.547	0.584474	
RegionSouthern Islands	418430.77	116502.87	3.592	0.000337	***
RegionTanglin	791742.64	89432.04	8.853	< 2e-16	***
RegionToa Payoh	-115697.50	72654.57	-1.592	0.111447	
Age	-36376.26	1980.19	-18.370	< 2e-16	***
PurchaserPrivate	17580.53	37595.88	0.468	0.640109	
Area:TenureLeasehold	-463.02	48.81	-9.487	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 697500 on 1980 degrees of freedom
Multiple R-squared: 0.8513, Adjusted R-squared: 0.8499
F-statistic: 596.7 on 19 and 1980 DF, p-value: < 2.2e-16

```
> pred3 <- predict(L3, newdata=cen[test,])
```

```
> mean((pred3-cen[test, "Price"])^2)
```

```
[1] 481090502793
```

```
> coef(L3)
```

(Intercept)	Area	TenureLeasehold
-40570.7492	2367.3548	233499.0720
RegionBukit Merah	RegionBukit Timah	RegionGeylang
206888.7958	77466.4604	-70344.6571
RegionKallang	RegionMarine Parade	RegionNewton
26012.7663	284265.3763	1357491.8971
RegionNovena	RegionOthers	RegionQueenstown
166281.2611	777659.3893	-49387.9679
RegionRiver Valley	RegionRochor	RegionSouthern Islands
975827.0232	70015.9712	418430.7747
RegionTanglin	RegionToa Payoh	Age
791742.6400	-115697.4961	-36376.2573
PurchaserPrivate	Area:TenureLeasehold	
17580.5330	-463.0165	

```
> #Examples of other lm models
```

```
> summary(lm(Price~.+I(Area^2)+I(Area^3), cen[train,]))
```

Call:

```
lm(formula = Price ~ . + I(Area^2) + I(Area^3), data = cen[train,
])
```

Residuals:

Min	1Q	Median	3Q	Max
-4976789	-159611	17575	192216	10129066

Coefficients:

Estimate	Std. Error	t value	Pr(> t)
----------	------------	---------	----------

(Intercept)	7.058e+05	1.148e+05	6.146	9.57e-10	***
Area	1.452e+03	1.383e+02	10.496	< 2e-16	***
Age	-3.388e+04	2.121e+03	-15.974	< 2e-16	***
TenureLeasehold	-3.589e+05	4.160e+04	-8.629	< 2e-16	***
PurchaserPrivate	6.962e+03	3.823e+04	0.182	0.85552	
RegionBukit Merah	1.192e+05	7.221e+04	1.651	0.09889	.
RegionBukit Timah	3.935e+04	7.182e+04	0.548	0.58381	
RegionGeylang	-9.376e+04	8.648e+04	-1.084	0.27840	
RegionKallang	-9.761e+03	7.981e+04	-0.122	0.90268	
RegionMarine Parade	3.006e+05	9.425e+04	3.189	0.00145	**
RegionNewton	1.428e+06	1.204e+05	11.857	< 2e-16	***
RegionNovena	1.221e+05	9.708e+04	1.257	0.20880	
RegionOthers	7.427e+05	1.165e+05	6.372	2.31e-10	***
RegionQueenstown	-6.246e+04	8.437e+04	-0.740	0.45921	
RegionRiver Valley	1.040e+06	1.041e+05	9.989	< 2e-16	***
RegionRochor	2.274e+04	1.304e+05	0.174	0.86160	
RegionSouthern Islands	1.243e+05	1.155e+05	1.076	0.28208	
RegionTanglin	8.196e+05	9.079e+04	9.028	< 2e-16	***
RegionToa Payoh	-9.068e+04	7.383e+04	-1.228	0.21951	
I(Area^2)	2.788e-01	5.409e-02	5.154	2.80e-07	***
I(Area^3)	-2.660e-05	5.530e-06	-4.811	1.62e-06	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 708400 on 1979 degrees of freedom

Multiple R-squared: 0.8467, Adjusted R-squared: 0.8452

F-statistic: 546.6 on 20 and 1979 DF, p-value: < 2.2e-16

```
> summary(lm(Price~Tenure*Purchaser+Region+Area+Age+I(Area^2), cen[train,]))
```

Call:

```
lm(formula = Price ~ Tenure * Purchaser + Region + Area + Age +
    I(Area^2), data = cen[train, ])
```

Residuals:

Min	1Q	Median	3Q	Max
-4751055	-174191	20108	196775	10485363

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	2.854e+05	1.037e+05	2.752	0.00597	**
TenureLeasehold	-2.446e+05	7.598e+04	-3.219	0.00131	**
PurchaserPrivate	1.091e+05	6.955e+04	1.569	0.11688	
RegionBukit Merah	1.477e+05	7.250e+04	2.038	0.04170	*
RegionBukit Timah	5.506e+04	7.215e+04	0.763	0.44545	
RegionGeylang	-7.188e+04	8.682e+04	-0.828	0.40779	
RegionKallang	1.667e+04	8.012e+04	0.208	0.83519	
RegionMarine Parade	2.893e+05	9.470e+04	3.055	0.00228	**
RegionNewton	1.428e+06	1.211e+05	11.786	< 2e-16	***
RegionNovena	1.399e+05	9.760e+04	1.433	0.15193	
RegionOthers	7.461e+05	1.171e+05	6.369	2.36e-10	***

RegionQueenstown	-2.149e+04	8.433e+04	-0.255	0.79888	
RegionRiver Valley	1.051e+06	1.046e+05	10.049	< 2e-16	***
RegionRochor	6.355e+04	1.309e+05	0.486	0.62733	
RegionSouthern Islands	1.725e+05	1.162e+05	1.485	0.13782	
RegionTanglin	8.413e+05	9.112e+04	9.233	< 2e-16	***
RegionToa Payoh	-6.699e+04	7.402e+04	-0.905	0.36553	
Area	2.050e+03	6.149e+01	33.336	< 2e-16	***
Age	-3.625e+04	2.076e+03	-17.464	< 2e-16	***
I (Area^2)	2.358e-02	1.106e-02	2.132	0.03314	*
TenureLeasehold:PurchaserPrivate	-1.508e+05	8.220e+04	-1.835	0.06668	.

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 711900 on 1979 degrees of freedom

Multiple R-squared: 0.8452, Adjusted R-squared: 0.8436

F-statistic: 540.3 on 20 and 1979 DF, p-value: < 2.2e-16

```
> summary(lm(Price~Tenure*Purchaser+Region+Area+Age, cen[train,]))
```

Call:

```
lm(formula = Price ~ Tenure * Purchaser + Region + Area + Age,
    data = cen[train, ])
```

Residuals:

	Min	1Q	Median	3Q	Max
	-4659550	-180644	24109	202650	10509849

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	194596.61	94628.45	2.056	0.03987	*
TenureLeasehold	-236626.80	75961.09	-3.115	0.00187	**
PurchaserPrivate	105668.47	69589.93	1.518	0.12906	
RegionBukit Merah	149887.88	72554.42	2.066	0.03897	*
RegionBukit Timah	57999.18	72201.55	0.803	0.42190	
RegionGeylang	-62616.77	86784.86	-0.722	0.47068	
RegionKallang	21985.28	80154.67	0.274	0.78389	
RegionMarine Parade	284172.02	94759.00	2.999	0.00274	**
RegionNewton	1407774.85	120886.85	11.645	< 2e-16	***
RegionNovena	141711.48	97686.89	1.451	0.14703	
RegionOthers	737737.18	117176.81	6.296	3.75e-10	***
RegionQueenstown	-6713.63	84116.83	-0.080	0.93639	
RegionRiver Valley	1054502.52	104700.21	10.072	< 2e-16	***
RegionRochor	79489.36	130792.46	0.608	0.54342	
RegionSouthern Islands	133005.29	114835.34	1.158	0.24691	
RegionTanglin	846677.89	91163.10	9.288	< 2e-16	***
RegionToa Payoh	-57367.21	73943.62	-0.776	0.43795	
Area	2168.61	25.91	83.709	< 2e-16	***
Age	-37274.15	2020.67	-18.446	< 2e-16	***
TenureLeasehold:PurchaserPrivate	-154727.89	82250.28	-1.881	0.06009	.

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 712600 on 1980 degrees of freedom
Multiple R-squared: 0.8448, Adjusted R-squared: 0.8434
F-statistic: 567.4 on 19 and 1980 DF, p-value: < 2.2e-16

Appendix 2 (Best Subset Selection)

Code:

```
set.seed(9876)
central <- read.csv("Central2024P.csv", stringsAsFactors = TRUE)
library(leaps)
regfit2 <- regsubsets(Price~., data=central, nvmax=18)
regfit2.summary <- summary(regfit2)
plot(regfit2.summary$adjr2, main="Adjusted r^2 plot", xlab="Number of
variables", ylab="Adjusted r^2", type="b")
plot(regfit2.summary$cp, main="Cp plot", xlab="Number of variables", ylab="cp",
type="b")
plot(regfit2.summary$bic, main="BIC plot", xlab="Number of variables",
ylab="BIC", type="b")
b <- which.max(regfit2.summary$adjr2)
c <- which.min(regfit2.summary$cp)
d <- which.min(regfit2.summary$bic)
# Model based on adjusted R square criteria
rsq <- coef(regfit2, b)
print(rsq)
# Model based on Cp criteria
cp <- coef(regfit2, c)
print(cp)
# Model based on BIC criteria
bic <- coef(regfit2, d)
print(bic)

# 10-fold cross validation on best subset
predict.regsubsets <- function(object, newdata, id){
  form <- as.formula(object$call[[2]])
  mat <- model.matrix(form, newdata)
  coefi <- coef(object, id=id)
  xvars <- names(coefi)
  mat[, xvars]%%coefi
}

k <- 10
set.seed(9876)
folds <- sample(1:k, nrow(central), replace=TRUE)
cv.errors <- matrix(NA, k, 18, dimnames=list(NULL, paste(1:18)))
for (j in 1:k) {
```

```

best.fit <- regsubsets(Price~., data=central[folds!=j,], nvmax=18)
for (i in 1:18){
  pred <- predict.regsubsets(best.fit, central[folds==j,], id=i)
  cv.errors[j,i] <- mean((central$Price[folds==j]-pred)^2)
}
}
# Test error
mean.cv <- apply(cv.errors, 2, mean)
min(mean.cv)
# Model with lowest cross validation error
bb <- which.min(mean.cv)
bssmod <- coef(regfit2, bb)
print(bssmod)

```

Code Output:

```

> set.seed(9876)
> central <- read.csv("Central2024P.csv", stringsAsFactors = TRUE)
> library(leaps)
> regfit2 <- regsubsets(Price~., data=central, nvmax=18)
> regfit2.summary <- summary(regfit2)
> plot(regfit2.summary$adjr2, main="Adjusted r^2 plot", xlab="Number of
variables", ylab="Adjusted r^2", type="b")
> plot(regfit2.summary$cp, main="Cp plot", xlab="Number of variables",
ylab="cp", type="b")
> plot(regfit2.summary$bic, main="BIC plot", xlab="Number of variables",
ylab="BIC", type="b")
> b <- which.max(regfit2.summary$adjr2)
> c <- which.min(regfit2.summary$cp)
> d <- which.min(regfit2.summary$bic)
> # Model based on adjusted R square criteria
> rsq <- coef(regfit2, b)
> print(rsq)

```

	(Intercept)	Area	Age
	304987.321	2169.022	-36864.128
TenureLeasehold	-354543.134	110757.155	-94923.726
RegionMarine Parade	253148.400	1565328.975	119923.266
RegionOthers	749686.525	935075.025	809621.314
RegionToa Payoh	-87135.277		
RegionBukit Merah			
RegionGeylang			
RegionNewton			
RegionNovena			
RegionTanglin			

```

> # Model based on Cp criteria
> cp <- coef(regfit2, c)
> print(cp)

```

	(Intercept)	Area	Age
	304987.321	2169.022	-36864.128
TenureLeasehold	-354543.134	110757.155	-94923.726
RegionMarine Parade			
RegionNewton			
RegionNovena			
RegionGeylang			
RegionTanglin			

```

      253148.400      1565328.975      119923.266
      RegionOthers RegionRiver Valley RegionTanglin
      749686.525      935075.025      809621.314
RegionToa Payoh
-87135.277
> # Model based on BIC criteria
> bic <- coef(regfit2, d)
> print(bic)
      (Intercept)      Area      Age
      288491.849      2173.528      -36113.552
TenureLeasehold RegionBukit Merah RegionMarine Parade
      -374889.486      130237.780      256243.812
      RegionNewton      RegionOthers RegionRiver Valley
      1563601.519      757482.451      939286.494
      RegionTanglin
      807644.255
>
> # 10-fold cross validation on best subset
> predict.regsubsets <- function(object, newdata, id){
+   form <- as.formula(object$call[[2]])
+   mat <- model.matrix(form, newdata)
+   coefi <- coef(object, id=id)
+   xvars <- names(coefi)
+   mat[, xvars] %*% coefi
+ }
>
> k <- 10
> set.seed(9876)
> folds <- sample(1:k, nrow(central), replace=TRUE)
> cv.errors <- matrix(NA, k, 18, dimnames=list(NULL, paste(1:18)))
> for (j in 1:k) {
+   best.fit <- regsubsets(Price~., data=central[folds!=j,], nvmax=18)
+   for (i in 1:18){
+     pred <- predict.regsubsets(best.fit, central[folds==j,], id=i)
+     cv.errors[j,i] <- mean((central$Price[folds==j]-pred)^2)
+   }
+ }
> # Test error
> mean.cv <- apply(cv.errors, 2, mean)
> min(mean.cv)
[1] 507174939759
> # Model with lowest cross validation error
> bb <- which.min(mean.cv)
> bssmod <- coef(regfit2, bb)
> print(bssmod)
      (Intercept)      Area      Age
      304987.321      2169.022      -36864.128
TenureLeasehold RegionBukit Merah RegionGeylang
      -354543.134      110757.155      -94923.726
RegionMarine Parade RegionNewton RegionNovena
      253148.400      1565328.975      119923.266
      RegionOthers RegionRiver Valley RegionTanglin
      749686.525      935075.025      809621.314

```

RegionToa Payoh
-87135.277

Appendix 3 (Ridge Regression)

Code:

```
set.seed(9876)
central <- read.csv("Central2024P.csv", stringsAsFactors = TRUE)
attach(central)
train <- sample(1:nrow(central), 2000)
test <- (-train)
library(glmnet)
x <- model.matrix(Price~., central)[, -1]
y <- central$Price
centraltrain <- central[train,]
centraltest <- central[test,]
trainx <- model.matrix(Price~., centraltrain)[, -1]
trainy <- centraltrain$Price
testx <- model.matrix(Price~., centraltest)[, -1]
testy <- centraltest$Price
ridgemod <- glmnet(trainx, trainy, alpha = 0)
cvout <- cv.glmnet(trainx, trainy, alpha = 0)
lambdarr <- cvout$lambda.min
lambdarr

# Calculating test error
ridgepred <- predict(ridgemod, s = lambdarr, newx = x[test,])
mean((ridgepred-testy)^2)

# Ridge regression model
outrr <- glmnet(x, y, alpha = 0)
rrmodel <- predict(outrr, type = "coefficients", s = lambdarr)[1:19,]
rrmodel[rrmodel!=0]
```

Code Output:

```
> set.seed(9876)
> central <- read.csv("Central2024P.csv", stringsAsFactors = TRUE)
> attach(central)
> train <- sample(1:nrow(central), 2000)
> test <- (-train)
> library(glmnet)
> x <- model.matrix(Price~., central)[, -1]
> y <- central$Price
> centraltrain <- central[train,]
```

```

> centraltest <- central[test,]
> trainx <- model.matrix(Price~., centraltrain)[, -1]
> trainy <- centraltrain$Price
> testx <- model.matrix(Price~., centraltest)[, -1]
> testy <- centraltest$Price
> ridgmod <- glmnet(trainx, trainy, alpha = 0)
> cvout <- cv.glmnet(trainx, trainy, alpha = 0)
> lambdarr <- cvout$lambda.min
> lambdarr
[1] 158285
>
> # Calculating test error
> ridgepred <- predict(ridgmod, s = lambdarr, newx = x[test,])
> mean((ridgepred-testy)^2)
[1] 5.19472e+11
>
> # Ridge regression model
> outrr <- glmnet(x, y, alpha = 0)
> rrmodel <- predict(outrr, type = "coefficients", s = lambdarr)[1:19,]
> rrmodel[rrmodel!=0]

```

(Intercept)	Area	Age
529570.255	1924.381	-27406.529
TenureLeasehold	PurchaserPrivate	RegionBukit Merah
-371811.706	53870.582	54012.431
RegionBukit Timah	RegionGeylang	RegionKallang
-75430.938	-188425.461	-84047.714
RegionMarine Parade	RegionNewton	RegionNovena
166561.947	1491826.575	27036.165
RegionOthers	RegionQueenstown	RegionRiver Valley
614129.229	-138935.274	858131.708
RegionRochor	RegionSouthern Islands	RegionTanglin
-123249.440	230595.753	694684.754
RegionToa Payoh		
-153123.336		

Appendix 4 (The Lasso)

Code:

```

set.seed(9876)
central <- read.csv("Central2024P.csv", stringsAsFactors = TRUE)
attach(central)
library(glmnet)
train <- sample(1:nrow(central), 2000)
test <- (-train)
x <- model.matrix(Price~., central)[, -1]
y <- central$Price

```



```

centraltrain <- central[train,]
centraltest <- central[test,]
trainx <- model.matrix(Price~., centraltrain)[, -1]
trainy <- centraltrain$Price
testx <- model.matrix(Price~., centraltest)[, -1]
testy <- centraltest$Price
lassomod <- glmnet(trainx, trainy, alpha = 1)
cvout1 <- cv.glmnet(trainx, trainy, alpha = 1)
lambdalasso <- cvout1$lambda.min
lambdalasso

# Test error
lassopred <- predict(lassomod, s = lambdalasso, newx = x[test,])
mean((lassopred-testy)^2)

# The lasso model
outlr <- glmnet(x, y, alpha = 1)
lrmodel <- predict(outlr, type = "coefficients", s = lambdalasso)[1:19,]
lrmodel[!lrmodel!=0]

```

Code Output:

```

> set.seed(9876)
> central <- read.csv("Central2024P.csv", stringsAsFactors = TRUE)
> attach(central)
> library(glmnet)
> train <- sample(1:nrow(central), 2000)
> test <- (-train)
> x <- model.matrix(Price~., central)[, -1]
> y <- central$Price
> centraltrain <- central[train,]
> centraltest <- central[test,]
> trainx <- model.matrix(Price~., centraltrain)[, -1]
> trainy <- centraltrain$Price
> testx <- model.matrix(Price~., centraltest)[, -1]
> testy <- centraltest$Price
> lassomod <- glmnet(trainx, trainy, alpha = 1)
> cvout1 <- cv.glmnet(trainx, trainy, alpha = 1)
> lambdalasso <- cvout1$lambda.min
> lambdalasso
[1] 1778.053
>
> # Test error
> lassopred <- predict(lassomod, s = lambdalasso, newx = x[test,])
> mean((lassopred-testy)^2)
[1] 506696419155
>
> # The lasso model
> outlr <- glmnet(x, y, alpha = 1)

```

```
> lrmodel <- predict(outlr, type = "coefficients", s = lambdalasso)[1:19,]
> lrmodel[lrmodel!=0]
```

(Intercept)	Area	Age
310677.535	2160.767	-36086.039
TenureLeasehold	RegionBukit Merah	RegionGeylang
-354979.350	98973.823	-89809.369
RegionMarine Parade	RegionNewton	RegionNovena
234580.741	1546302.832	100410.572
RegionOthers	RegionQueenstown	RegionRiver Valley
725095.604	-19342.087	917529.075
RegionRochor	RegionSouthern Islands	RegionTanglin
-9852.625	44343.389	791096.014
RegionToa Payoh		
-81996.751		

Appendix 5 (Elastic Net Regression)

The code below is run using alpha=0.25

Code:

```
central2024 <- read.csv("Central2024P.csv", stringsAsFactors = TRUE)
attach(central2024)
library(glmnet)
set.seed(9876)
train_index <- sample(1:nrow(central2024), 2000)
train_data <- central2024[train_index, ]
test_data <- central2024[-train_index, ]
x.train <- model.matrix(Price ~ Area + Age + Tenure + Purchaser + Region, data =
train_data)[, -1]
y.train <- train_data$Price

#set alpha = 0.5 for Elastic Net (0 for Ridge, 1 for Lasso)
elasticnet <- cv.glmnet(x.train, y.train, alpha = 0.25)

#cross validation results
plot(elasticnet)
bestlambda <- elasticnet$lambda.min
final_model <- glmnet(x.train, y.train, alpha = 0.5, lambda = bestlambda)

# use model for testing
x.test <- model.matrix(Price ~ Area + Age + Tenure + Purchaser + Region, data =
test_data)[, -1]
y.pred <- predict(final_model, newx = x.test)
MSE <- mean((test_data$Price - y.pred)^2)
MSE
bestlambda
model_coefficients <- coef(final_model)
print(model_coefficients)
```

Code output:

```
> central2024 <- read.csv("Central2024P.csv", stringsAsFactors = TRUE)
> attach(central2024)
> library(glmnet)
Loading required package: Matrix
Loaded glmnet 4.1-8
> set.seed(9876)
> train_index <- sample(1:nrow(central2024), 2000)
> train_data <- central2024[train_index, ]
> test_data <- central2024[-train_index, ]
> x.train <- model.matrix(Price ~ Area + Age + Tenure + Purchaser + Region,
data = train_data)[, -1]
> y.train <- train_data$Price

> #set alpha = 0.5 for Elastic Net (0 for Ridge, 1 for Lasso)
> elasticnet <- cv.glmnet(x.train, y.train, alpha = 0.25)

> #cross validation results
> plot(elasticnet)
> bestlambda <- elasticnet$lambda.min
> final_model <- glmnet(x.train, y.train, alpha = 0.5, lambda = bestlambda)

> # use model for testing
> x.test <- model.matrix(Price ~ Area + Age + Tenure + Purchaser + Region,
data = test_data)[, -1]
> y.pred <- predict(final_model, newx = x.test)
> MSE <- mean((test_data$Price - y.pred)^2)
> MSE
[1] 506447279753
> bestlambda
[1] 5380.128
> model_coefficients <- coef(final_model)
> print(model_coefficients)
19 x 1 sparse Matrix of class "dgCMatrix"
               s0
(Intercept)    3.250865e+05
Area           2.163697e+03
Age            -3.642627e+04
TenureLeasehold -3.606519e+05
PurchaserPrivate .
RegionBukit Merah    9.293030e+04
RegionBukit Timah    2.464985e+00
RegionGeylang       -9.782176e+04
RegionKallang        -2.124992e+04
RegionMarine Parade  2.243633e+05
RegionNewton         1.354373e+06
RegionNovena         7.458192e+04
RegionOthers         6.692799e+05
RegionQueenstown    -4.151963e+04
RegionRiver Valley   9.969955e+05
RegionRochor         1.089719e+04
RegionSouthern Islands 6.277788e+04
```

RegionTanglin	7.852705e+05
RegionToa Payoh	-8.718872e+04

Appendix 6 (K-nearest Neighbours)

Code:

```
x <- read.csv("Central2024P.csv", stringsAsFactors = TRUE)
View(x)
library(class)

x$Tenure <- as.numeric(x$Tenure)
x$Purchaser <- as.numeric(x$Purchaser)
x$Region <- as.numeric(x$Region)
#x[,2:6] <- scale(x[,2:6])
View(x)
set.seed(9876)
num_folds <- 5

# Vector to store the mean squared errors for each fold
mse_cv <- numeric(num_folds)

# Perform k-fold cross-validation
for (i in 1:num_folds) {
  # Define the indices for the current fold
  fold_indices <- sample(1:nrow(x), size = nrow(x) / num_folds)

  # Split the data into training and validation sets
  train_fold <- x[-fold_indices, ]
  validation_fold <- x[fold_indices, ]

  # KNN model
  knn_model <- knn(train = train_fold[, -which(names(train_fold) == "Price")],
                   test = validation_fold[, -which(names(validation_fold) ==
"Price")],
                   cl = train_fold$Price,
                   k = 4) # You can adjust the value of k as needed

  predictions <- as.numeric(knn_model)

  mse_cv[i] <- mean((predictions - validation_fold$Price)^2)
}

# Calculate the average Mean Squared Error across all folds
average_mse_cv <- mean(mse_cv)

print(paste("Average Mean Squared Error (Cross-Validation):", average_mse_cv))
```

Code Output:

```
> x <- read.csv("Central2024P.csv", stringsAsFactors = TRUE)
> library(class)
> x$Tenure <- as.numeric(x$Tenure)
> x$Purchaser <- as.numeric(x$Purchaser)
> x$Region <- as.numeric(x$Region)
> set.seed(9876)
> num_folds <- 5
> # Vector to store the mean squared errors for each fold
> mse_cv <- numeric(num_folds)
> # Perform k-fold cross-validation
> for (i in 1:num_folds) {
+   # Define the indices for the current fold
+   fold_indices <- sample(1:nrow(x), size = nrow(x) / num_folds)
+
+   # Split the data into training and validation sets
+   train_fold <- x[-fold_indices, ]
+   validation_fold <- x[fold_indices, ]
+
+   # KNN model
+   knn_model <- knn(train = train_fold[, -which(names(train_fold) ==
"Price")],
+                     test = validation_fold[, -which(names(validation_fold) ==
"Price")],
+                     cl = train_fold$Price,
+                     k = 4) # You can adjust the value of k as needed
+
+   predictions <- as.numeric(knn_model)
+
+   mse_cv[i] <- mean((predictions - validation_fold$Price)^2)
+ }
> # Calculate the average Mean Squared Error across all folds
> average_mse_cv <- mean(mse_cv)
> print(paste("Average Mean Squared Error (Cross-Validation):",
average_mse_cv))
[1] "Average Mean Squared Error (Cross-Validation): 9285272687764.77"
```

Appendix 7 (Decision Tree)

Code:

```
#Decision Tree Approach
```

```
x <- read.csv("Central2024P.csv", stringsAsFactors = TRUE)
View(x)
#x[,1] <- scale(x[,1])
#attach(x)
library(tree)
```

```

set.seed(9876)
train <- sample(1:nrow(x),2000)
test <- -train

#Initial Tree construction
tree.central <- tree(Price~., data = x, subset = train)
summary(tree.central)
tree.central
plot(tree.central)
title("Regression Tree for Central2024P data")
text(tree.central, pretty = 0, cex = 0.6, srt = 5)

#Finding minimum nodes through cv

cv.central <- cv.tree(tree.central)
cv.central
plot(cv.central$size, cv.central$dev, type="b", main="Cross validation: Deviance
versus Size",
      xlab="Number of terminal nodes", ylab="deviance")
minimum_nodes <- cv.central$size[which.min(cv.central$dev)]
minimum_nodes

#Pruning given minimum nodes

prune.central <- prune.tree(tree.central, best=minimum_nodes)
plot(prune.central)
title ("Pruned Regression Tree for central data")
text(prune.central, pretty=0, cex = 0.6, srt = 5)

predict <- predict(prune.central, newdata = x[test,])
central.test <- x[test, 'Price']

plot(predict, central.test, main="Pruned Tree prediction versus observed prices
for test data",
      xlab="predict Price", ylab="Observed Price")
mean((predict-central.test)^2)

```

Code Output:

```

> x <- read.csv("Central2024P.csv", stringsAsFactors = TRUE)
> View(x)
> #x[,1] <- scale(x[,1])
> #attach(x)
> library(tree)
> set.seed(9876)
> train <- sample(1:nrow(x),2000)
> test <- -train
> #Initial Tree construction
> tree.central <- tree(Price~., data = x, subset = train)
> summary(tree.central)

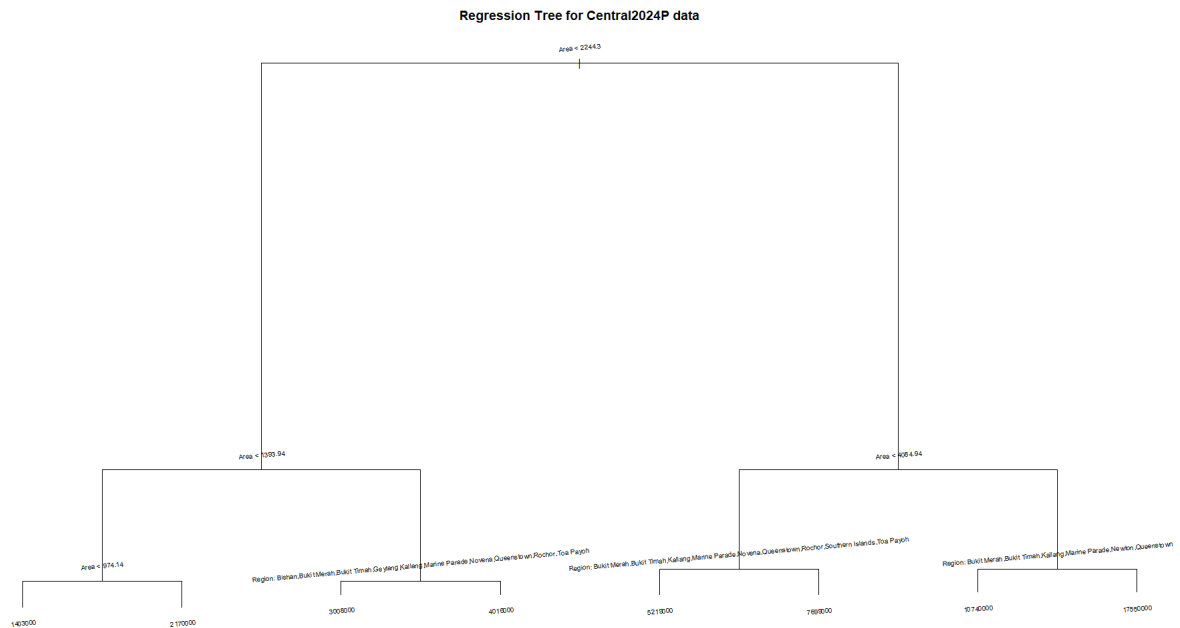
```

```

Regression tree:
tree(formula = Price ~ ., data = x, subset = train)
Variables actually used in tree construction:
[1] "Area"    "Region"
Number of terminal nodes: 8
Residual mean deviance: 5.499e+11 = 1.095e+15 / 1992
Distribution of residuals:
      Min. 1st Qu.  Median      Mean 3rd Qu.    Max.
-3919000 -353900  -36990         0  277200 10300000
> tree.central
node), split, n, deviance, yval
      * denotes terminal node

1) root 2000 6.480e+15 2522000
  2) Area < 2244.3 1841 1.611e+15 2156000
    4) Area < 1393.94 1368 4.066e+14 1757000
      8) Area < 974.14 737 9.241e+13 1403000 *
      9) Area > 974.14 631 1.139e+14 2170000 *
    5) Area > 1393.94 473 3.557e+14 3310000
      10) Region: Bishan,Bukit Merah,Bukit Timah,Geylang,Kallang,Marine
Parade,Novena,Queenstown,Rochor,Toa Payoh 331 1.195e+14 3008000 *
      11) Region: Newton,Others,River Valley,Southern Islands,Tanglin 142
1.352e+14 4016000 *
    3) Area > 2244.3 159 1.760e+15 6765000
      6) Area < 4084.94 143 6.956e+14 6034000
        12) Region: Bukit Merah,Bukit Timah,Kallang,Marine
Parade,Novena,Queenstown,Rochor,Southern Islands,Toa Payoh 96 1.800e+14
5219000 *
        13) Region: Newton,Others,River Valley,Tanglin 47 3.218e+14 7699000 *
      7) Area > 4084.94 16 3.063e+14 13290000
        14) Region: Bukit Merah,Bukit Timah,Kallang,Marine
Parade,Newton,Queenstown 10 8.227e+13 10740000 *
        15) Region: River Valley,Tanglin 6 5.047e+13 17550000 *
> plot(tree.central)
> title("Regression Tree for Central2024P data")
> text(tree.central, pretty = 0, cex = 0.6, srt = 5)

```



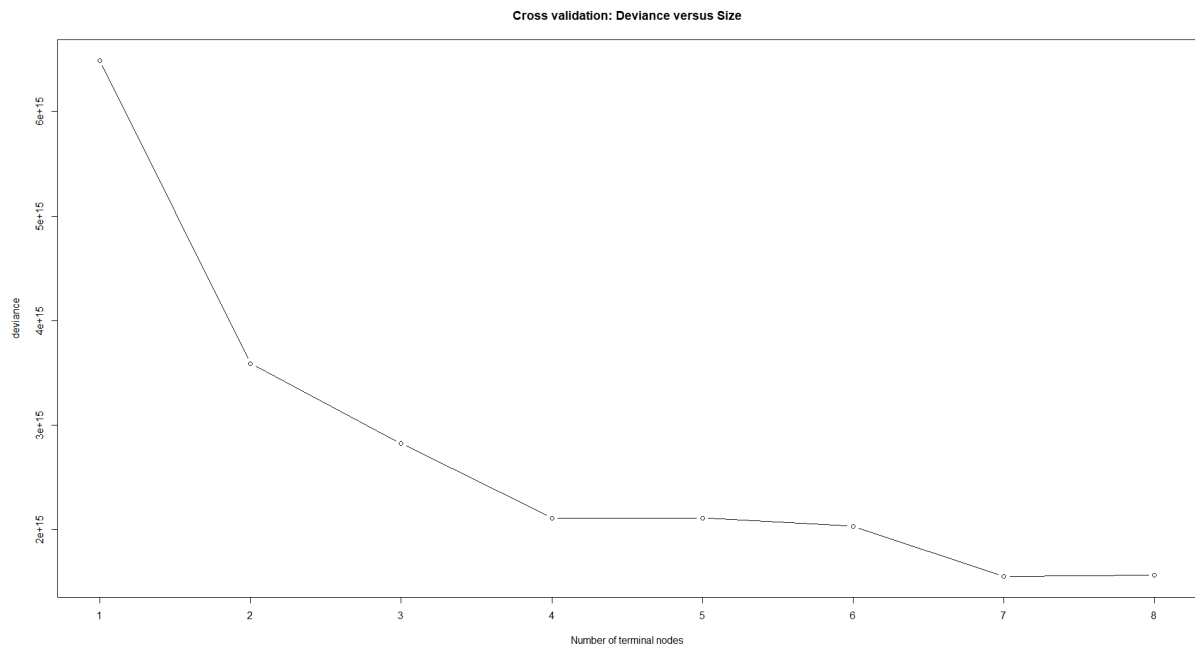
```
> cv.central <- cv.tree(tree.central)
> cv.central
$size
[1] 8 7 6 5 4 3 2 1

$dev
[1] 1.565915e+15 1.552796e+15 2.033781e+15 2.111384e+15 2.111384e+15
2.823611e+15 3.588331e+15
[8] 6.487844e+15

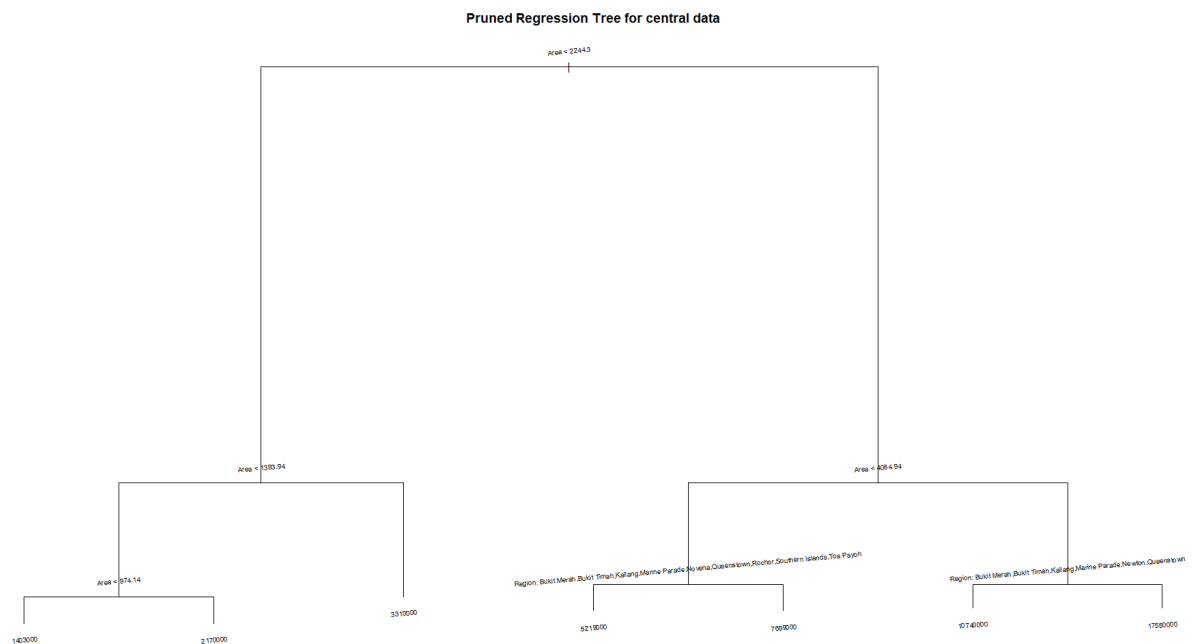
$k
[1] -Inf 1.010539e+14 1.735927e+14 1.939137e+14 2.003051e+14
7.584704e+14 8.483383e+14
[8] 3.108814e+15

$method
[1] "deviance"

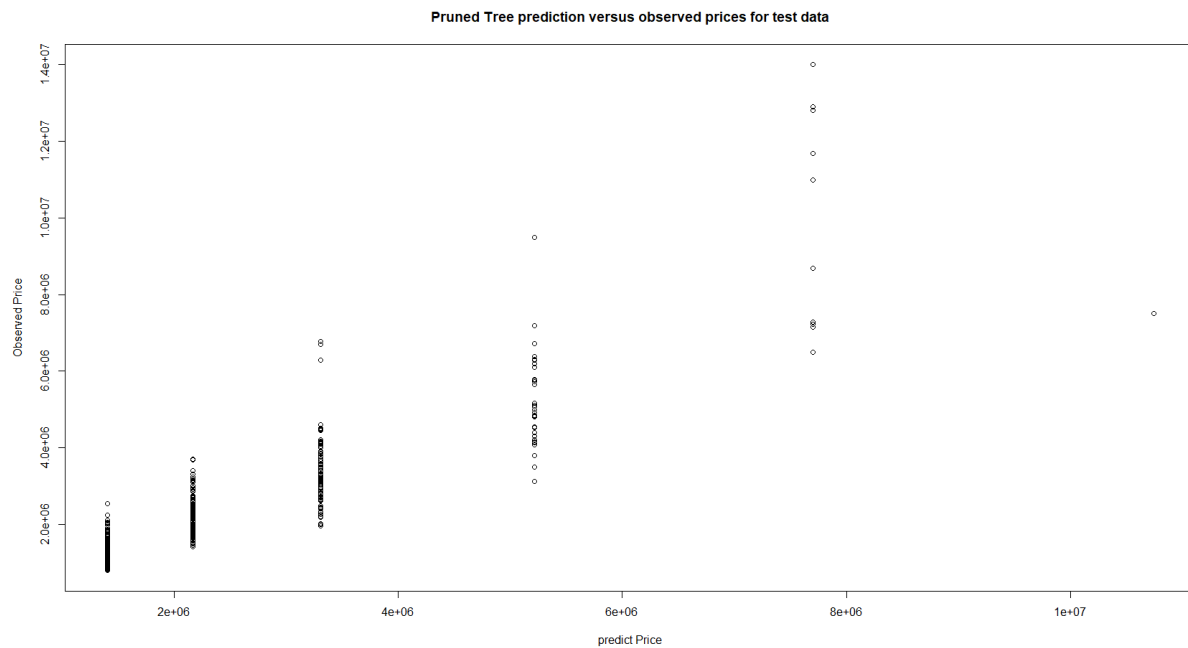
attr(,"class")
[1] "prune" "tree.sequence"
> plot(cv.central$size, cv.central$dev, type="b", main="Cross validation:
Deviance versus Size",
+ xlab="Number of terminal nodes", ylab="deviance")
> minimum_nodes <- cv.central$size[which.min(cv.central$dev)]
> minimum_nodes
[1] 7
```

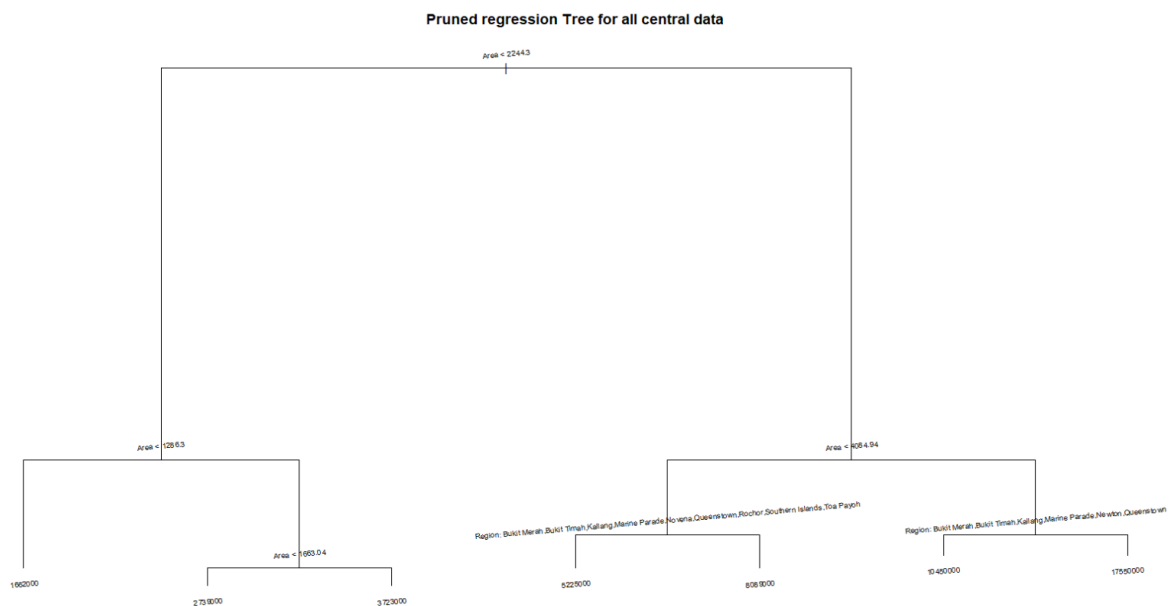
```
> prune.central <- prune.tree(tree.central, best=minimum_nodes)
> plot(prune.central)
> title ("Pruned Regression Tree for central data")
> text(prune.central, pretty=0, cex = 0.6, srt = 5)
```



```
> predict <- predict(prune.central, newdata = x[test,])
> central.test <- x[test, 'Price']
> plot(predict, central.test, main="Pruned Tree prediction versus observed
prices for test data",
+       xlab="predict Price", ylab="Observed Price")
> mean((predict-central.test)^2)
[1] 6.30826e+11
```



```
> #Build model with all data
> tree.centralall <- tree(Price~., data = x)
> prune.centralall <- prune.tree(tree.centralall, best = minimum_nodes)
> plot(prune.centralall)
> title("Pruned regression Tree for all central data")
> text(prune.centralall, pretty=0, cex = 0.6, srt = 5)
```



```
#Testing
> predict2 <- predict(prune.centralall, newdata = x[test,])
> central.test <- x[test, 'Price']
> mean((predict2-central.test)^2)
[1] 625661162260
```

Appendix 8 (Random Forest)

Code:

```
x <- read.csv("Central2024P.csv", stringsAsFactors = TRUE)
library(tree)
set.seed(9876)
train <- sample(1:nrow(x), 2000)
test <- -train

#Attempting random forest
library(randomForest)
predictor_variables <- names(x)[-which(names(x) == "Price")]
target_variable <- "Price"

# Train the Random Forest model using only the training data
rf_model <- randomForest(
  formula = as.formula(paste(target_variable, "~ .")),
  data = x[train,], # Use only the training data
  ntree = 550, # Number of trees in the forest
  mtry = 4, # Number of variables randomly sampled as candidates at each split
  importance = TRUE # Calculate variable importance
)

# Print the model summary
print(rf_model)

# Predict on the test data
predictions <- predict(rf_model, newdata = x[test,])

# Calculate Mean Squared Error
mse <- mean((predictions - x[test,]$Price)^2)
print(paste("Mean Squared Error:", mse))

# Get variable importance measures
importance <- importance(rf_model)
print(importance)

# Plot variable importance
varImpPlot(rf_model)
```

Output Code:

```
> x <- read.csv("Central2024P.csv", stringsAsFactors = TRUE)
> #attach(x)
> library(tree)
> set.seed(9876)
```

```

> train <- sample(1:nrow(x),2000)
> test <- -train
> #Attempting random forest
> library(randomForest)
> predictor_variables <- names(x)[-which(names(x) == "Price")]
> target_variable <- "Price"
> # Train the Random Forest model using only the training data
> rf_model <- randomForest(
+ formula = as.formula(paste(target_variable, "~ .")),
+ data = x[train,], # Use only the training data
+ ntree = 550, # Number of trees in the forest
+ mtry = 4, # Number of variables randomly sampled as candidates at each
split
+ importance = TRUE # Calculate variable importance
+ )
> # Print the model summary
> print(rf_model)

Call:
randomForest(formula = as.formula(paste(target_variable, "~ .")), data
= x[train, ], ntree = 550, mtry = 4, importance = TRUE)
      Type of random forest: regression
      Number of trees: 550
No. of variables tried at each split: 4

      Mean of squared residuals: 335182286517
      % Var explained: 89.65
> # Predict on the test data
> predictions <- predict(rf_model, newdata = x[test,])
> # Calculate Mean Squared Error
> mse <- mean((predictions - x[test,]$Price)^2)
> print(paste("Mean Squared Error:", mse))
[1] "Mean Squared Error: 221844798581.846"
> # Get variable importance measures
> importance <- importance(rf_model)
> print(importance)
      %IncMSE IncNodePurity
Area      138.417424  5.082534e+15
Age       39.320319  3.205572e+14
Tenure    23.649001  6.350828e+13
Purchaser  5.333474  6.472522e+12
Region    49.823769  8.945840e+14
> # Plot variable importance
> varImpPlot(rf_model)

```

rf_model

